

# Appendix to “Log-Structured Merge Design of Authenticated Data Structures for Efficiently Verifiable Cloud Outsourcing”

## A. CONSISTENCY CHECKING

```

1 class store_wrapper{
2   Store store;
3   Att cPut(key,val){
4     prePut(key,val);
5     att(tsw)=store.dPut(key,val);
6     return postPut(key,val,att(tsw));
7   }
8   Crt cGet(key){
9     preGet(<key>);
10    <key,val>,pf(tsrw,tsr*)=store.dGet(key);
11    return postGet(<key>,pf(tsrw,tsr*));
12  }
13
14  mutex State pending_wr, completed_wr,history_w;
15
16  void prePut(<key,val>){
17    pending_wr.add(<key,val,start_rt=now()>);
18  }
19  boolean postPut(<key,val>,att(tsw)){
20    <key,val,start_rt>=pending_wr.remove();
21    completed_wr.addW(<key,val,start_rt,end_rt=now(),tsw>);
22    ;
23    acl = completed_wr.tryTruncate();
24    if(acl != NULL){
25      assertC(acl,history_w);
26      for(Write w in acl.trim())
27        history_w.put(w);
28    }
29  }
30  void preGet(key){
31    pending_wr.add(<key,start_rt=now()>);
32  }
33  boolean postGet(r<key,val,tsrw,tsr,pf(tsrw,tsr*)>){
34    r<key,start_rt>=pending_wr.remove();
35    if(r.tsr <= history_w.latest()){
36      assertL2(r<key,val,tsr,tsrw,pf(tsrw,tsr*)>,history_w);
37    }
38    else
39      completed_wr.addR(<key,val,start_rt,end_rt=now(),tsr,tsrw>);
40  }
41 }

```

Listing 2: Interfaces of verified and verifiable Put/Get

```

1 void assertC(acl,history_w){
2   //L1
3   o0<key,val,ts> = history_w.latest();
4   o1<key,val,ts> = acl.oldest();
5   do {
6     assertL1pairwise(o0<key,val,ts>,o1<key,val,ts>);
7     o0=o1; o1=o1.next(acl);
8   } while (o1 != NULL)
9   //L2
10  for (read r in acl)
11    assertL2(r<key,val,tsr,tsrw,pf(tsrw,tsr*)>,history_w);
12 }
13 void assertL2(r<key,val,tsr,tsrw,pf(tsrw,tsr*)>,history_w)
14 {
15   assert(r.tsr <= history_w.latest());
16   assert(verify(pf(tsrw,tsr*))==true);
17   assert(tsr==tsr);
18 }

```

Listing 3: Linearizability checking

## B. FORMAL FRONTEND SECURITY

Intuitively, if a record is fresh in a set of records, the record must be fresh (or absent) in any subset of the set. The correctness comes from the intuition that “a record is the freshest in a dataset iff. it

is the freshest or absent in the subsets whose union fully covers the dataset.” And the unforgeability comes from the collision-resistance of the hashes and Merkle trees. *F1* can be proven with unforgeability by the authentication path of MHT digesting  $C_1$  and the key ordering among records in  $C_1$  implied Invariant 3.1. *F2* is proven with unforgeability by the authentication path at  $C_0$  and Invariant 3.1. *F3* is implied by invariant 3.3. Generalizing this case gives us a formal proof described below:

PROOF. Consider  $\langle k, v, t \rangle$  resides in level  $C_i$  of an LSM of  $n$  levels,  $C_j$  where  $j \in [1, n]$ .

Record  $\langle k, v, t \rangle$  is the freshest (with the largest timestamp value) among all records of key  $k$  in the entire LSM tree, iff the following four facts hold:

- F1 No record of key  $k$  exists in levels  $C_0, C_2, \dots, C_{i-1}$ .
- F2 Record  $\langle k, v, t \rangle @ C_i$  is the freshest among all records of key  $k$  in level  $C_i$ .
- F3 Records of key  $k$  in levels  $C_{i+1}, \dots, C_n$ , if they exist, are older than record  $\langle k, v, t \rangle$ .
- F4  $C_0, C_1, \dots, C_n$  covers the entire dataset in the LSM tree.

The four facts can be proved by the following proofs, respectively:

F1 is about the non-membership of records  $\langle k, \cdot \rangle$  on levels  $C_0, \dots, C_n$ . F1 can be verified by the authentication paths of key  $k$  of MHTs on these levels.

F2 can be similarly proved by the authentication path of key  $k$  of the MHT on level  $C_i$ .

F3 can be derived from the inter-level time-ordering (Invariant 3.3), as any record of key  $k$  on levels  $C_{>i}$  must have their timestamp smaller (older) than any record on level  $C_i$ .

F4 can be proved by that the enclave maintains the complete set of per-level Merkle hashes on all levels.  $\square$